

microsoft-stock-data

August 11, 2023

Microsoft Stock Data

```
[ ]: #import Library

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import os
import time
```

```
[ ]: #Load DataSet
from google.colab import files
upload=files.upload()
```

<IPython.core.display.HTML object>

Saving MSFT.csv to MSFT.csv

```
[ ]: df=pd.read_csv('MSFT.csv')
```

```
[ ]: df.head()
```

```
[ ]:
```

	Date	Open	High	Low	Close	Adj Close	Volume
0	1986-03-13	0.088542	0.101563	0.088542	0.097222	0.061434	1031788800
1	1986-03-14	0.097222	0.102431	0.097222	0.100694	0.063628	308160000
2	1986-03-17	0.100694	0.103299	0.100694	0.102431	0.064725	133171200
3	1986-03-18	0.102431	0.103299	0.098958	0.099826	0.063079	67766400
4	1986-03-19	0.099826	0.100694	0.097222	0.098090	0.061982	47894400

```
[ ]: df.info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 9083 entries, 0 to 9082

Data columns (total 7 columns):

#	Column	Non-Null Count	Dtype
0	Date	9083 non-null	object
1	Open	9083 non-null	float64

```

2   High      9083 non-null   float64
3   Low       9083 non-null   float64
4   Close     9083 non-null   float64
5   Adj Close 9083 non-null   float64
6   Volume    9083 non-null   int64
dtypes: float64(5), int64(1), object(1)
memory usage: 496.9+ KB

```

```
[ ]: df.set_index('Date',drop=True,inplace=True)
```

```
[ ]: df.head()
```

```
[ ]:
```

	Open	High	Low	Close	Adj Close	Volume
Date						
1986-03-13	0.088542	0.101563	0.088542	0.097222	0.061434	1031788800
1986-03-14	0.097222	0.102431	0.097222	0.100694	0.063628	308160000
1986-03-17	0.100694	0.103299	0.100694	0.102431	0.064725	133171200
1986-03-18	0.102431	0.103299	0.098958	0.099826	0.063079	67766400
1986-03-19	0.099826	0.100694	0.097222	0.098090	0.061982	47894400

We'll use Only Close Features

```
[ ]: df=df[['Close']]
```

```
[ ]: df
```

```
[ ]:
```

	Close
Date	
1986-03-13	0.097222
1986-03-14	0.100694
1986-03-17	0.102431
1986-03-18	0.099826
1986-03-19	0.098090
...	...
2022-03-18	300.429993
2022-03-21	299.160004
2022-03-22	304.059998
2022-03-23	299.489990
2022-03-24	304.100006

[9083 rows x 1 columns]


```
[ ]: df.describe()
```

```
[ ]:
```

	Close
count	9083.000000
mean	41.335628
std	59.714567

```
min      0.090278
25%      4.075195
50%     26.840000
75%     39.937500
max     343.109985
```

```
[ ]: plt.figure (1,figsize=(12,5))
plt.plot(df.Close);
```



Calculate the percentage Change **The Reason for using pct_change Instead of the prices is the beniefit of normalization as we can measure all variables a comparable metric. Also returns have more manageable statistics properties than prices such as stationarity, as in most cases we don't have stationary prices but we can have stationary returns.

A Stationary time series is one where statistical properties such as mean ,variance ,correlation ,etc are constant over Time**

```
[ ]: df['returns']=df.Close.pct_change()
```

```
[ ]: 134.75-132.89-1
```

```
[ ]: 0.860000000000000136
```

Calculate the log returns

***Several benefits of using log returns, both theoretic and algorithmic.

First, log-normality: if we assume that prices are distributed log normally (which, in practice, may or may not be true for any given price series), then $\log(1 + r_i)$ is conveniently normally distributed, because:


$$1 + r_i = \frac{p_i}{p_j} = \exp^{\{\log(\frac{p_i}{p_j})\}}$$

This is handy given much of classic statistics presumes normality***

```
[ ]: df['log_returns']=np.log(1+df['returns'])
```

```
[ ]: plt.figure(1,figsize=(16,6))  
plt.plot(df.log_returns)
```

```
[ ]: [ <matplotlib.lines.Line2D at 0x7bd5d531b220>]
```



```
[36]: #Drop NULL  
df.dropna(inplace=True)  
x=df[['Close','log_returns']].values
```

```
[37]: x
```

```
[37]: array([[ 1.00694000e-01,  3.50891917e-02],  
          [ 1.02431000e-01,  1.71031861e-02],  
          [ 9.98260000e-02, -2.57607307e-02],  
          ...,  
          [ 3.04059998e+02,  1.62464831e-02],  
          [ 2.99489990e+02, -1.51440491e-02],  
          [ 3.04100006e+02,  1.52756198e-02]])
```

```
[52]: #Scaling  
from sklearn.preprocessing import StandardScaler
```

```
[54]: scaler=StandardScaler()  
x_scaler=scaler.fit_transform(x)
```

```
[55]: x_scaler[:5]
```

```
[55]: array([[ -0.69062809,  1.59840864],  
          [ -0.690599  ,  0.75786948],  
          [ -0.69064263, -1.24528795],
```

```
[-0.6906717 , -0.86126061],  
[-0.69071531, -1.2987997 ]])
```

```
[95]: y=[x[0] for x in x_scaler ]
```

```
[101]: y[:5]
```

```
[101]: [-0.6906280902266285,  
-0.6905990010829361,  
-0.6906426264250884,  
-0.6906716988220084,  
-0.6907153074173882]
```

```
[58]: #Train Test Split
```

```
[97]: split=int(len(x_scaler)*0.8)  
print(split)
```

```
7265
```

```
[98]: x_train=x_scaler[:split]  
x_test=x_scaler[split: len(x_scaler)]  
y_train=y[:split]  
y_test=y[split: len(y)]
```

```
[100]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x_scaler,y,train_size=0.  
↪75,random_state=0)
```

```
[102]: len(x_train)==len(y_train)  
len(x_test)==len(y_test)
```

```
[102]: True
```

Labeling *we want to pridict the stock price at a future time .

as we going to use LSTM*

```
[161]: n=3  
xtrain=[]  
ytrain=[]  
xtest=[]  
ytest=[]  
for i in range(n,len(x_train)):  
    xtrain.append(x_train[i-n:i,:x_train.shape[1]]) #pridiction next record  
    ytrain.append(y_train[i])  
for i in range(n,len(x_test)):  
    ytest.append(y_test[i])
```

```
xtest.append(x_test[i-n:i,:x_test.shape[1]]) #prediction next record
```

```
[162]: df.head()
```

```
[162]:
```

	Close	returns	log_returns
Date			
1986-03-14	0.100694	0.035712	0.035089
1986-03-17	0.102431	0.017250	0.017103
1986-03-18	0.099826	-0.025432	-0.025761
1986-03-19	0.098090	-0.017390	-0.017543
1986-03-20	0.095486	-0.026547	-0.026906

```
[163]: xtrain[0]
```

```
[163]: array([[ -0.21972049,  0.13279207],  
          [ -0.49710982,  0.74827861],  
          [ 1.62074973,  0.58533444]])
```

```
[164]: ytrain[0]
```

```
[164]: -0.28310702240970176
```

```
[165]: val=np.array(ytrain[0])  
val=np.c_[val,np.zeros(val.shape)]
```

```
[166]: scaler.inverse_transform(val)
```

```
[166]: array([[2.44349990e+01, 8.86161074e-04]])
```

```
[167]: xtrain,ytrain=[np.array(xtrain),np.array(ytrain)]
```

```
[168]: xtrain=np.reshape(xtrain ,(xtrain.shape[0],xtrain.shape[1],xtrain.shape[2]))  
xtest=np.array(xtest)  
ytest=np.array(ytest)  
xtest=np.reshape(xtest ,(xtest.shape[0],xtest.shape[1],xtest.shape[2]))
```

```
[169]: print(xtrain.shape)  
print(ytrain.shape)  
print("*****")  
print(xtest.shape)  
print(ytest.shape)
```

```
(6808, 3, 2)
```

```
(6808,)
```

```
*****
```

```
(2268, 3, 2)
```

```
(2268,)
```

```
[170]: ##LSTM Model
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM,Dense

[171]: model=Sequential()

[172]: model.add(LSTM(4,input_shape=(xtrain.shape[1],xtrain.shape[2])))
model.add(Dense(1))
model.compile(loss='mean_squared_error',optimizer='adam')

[173]: model.
↳fit(xtrain,ytrain,validation_data=(xtest,ytest),epochs=50,batch_size=16,verbose=1)
```

```
Epoch 1/50
426/426 [=====] - 5s 5ms/step - loss: 1.0259 -
val_loss: 0.9458
Epoch 2/50
426/426 [=====] - 2s 4ms/step - loss: 1.0205 -
val_loss: 0.9433
Epoch 3/50
426/426 [=====] - 2s 5ms/step - loss: 1.0196 -
val_loss: 0.9428
Epoch 4/50
426/426 [=====] - 2s 4ms/step - loss: 1.0195 -
val_loss: 0.9430
Epoch 5/50
426/426 [=====] - 2s 5ms/step - loss: 1.0194 -
val_loss: 0.9424
Epoch 6/50
426/426 [=====] - 2s 5ms/step - loss: 1.0191 -
val_loss: 0.9433
Epoch 7/50
426/426 [=====] - 2s 4ms/step - loss: 1.0195 -
val_loss: 0.9424
Epoch 8/50
426/426 [=====] - 2s 4ms/step - loss: 1.0191 -
val_loss: 0.9421
Epoch 9/50
426/426 [=====] - 2s 4ms/step - loss: 1.0192 -
val_loss: 0.9420
Epoch 10/50
426/426 [=====] - 2s 4ms/step - loss: 1.0191 -
val_loss: 0.9420
Epoch 11/50
426/426 [=====] - 2s 4ms/step - loss: 1.0189 -
val_loss: 0.9420
Epoch 12/50
```

426/426 [=====] - 2s 5ms/step - loss: 1.0189 -
val_loss: 0.9419
Epoch 13/50
426/426 [=====] - 2s 5ms/step - loss: 1.0190 -
val_loss: 0.9417
Epoch 14/50
426/426 [=====] - 2s 4ms/step - loss: 1.0188 -
val_loss: 0.9423
Epoch 15/50
426/426 [=====] - 2s 4ms/step - loss: 1.0189 -
val_loss: 0.9418
Epoch 16/50
426/426 [=====] - 2s 4ms/step - loss: 1.0189 -
val_loss: 0.9415
Epoch 17/50
426/426 [=====] - 3s 6ms/step - loss: 1.0188 -
val_loss: 0.9420
Epoch 18/50
426/426 [=====] - 4s 9ms/step - loss: 1.0189 -
val_loss: 0.9417
Epoch 19/50
426/426 [=====] - 3s 6ms/step - loss: 1.0187 -
val_loss: 0.9419
Epoch 20/50
426/426 [=====] - 3s 6ms/step - loss: 1.0187 -
val_loss: 0.9414
Epoch 21/50
426/426 [=====] - 3s 6ms/step - loss: 1.0188 -
val_loss: 0.9413
Epoch 22/50
426/426 [=====] - 3s 7ms/step - loss: 1.0187 -
val_loss: 0.9413
Epoch 23/50
426/426 [=====] - 3s 7ms/step - loss: 1.0186 -
val_loss: 0.9414
Epoch 24/50
426/426 [=====] - 2s 5ms/step - loss: 1.0184 -
val_loss: 0.9411
Epoch 25/50
426/426 [=====] - 2s 6ms/step - loss: 1.0184 -
val_loss: 0.9416
Epoch 26/50
426/426 [=====] - 2s 6ms/step - loss: 1.0185 -
val_loss: 0.9412
Epoch 27/50
426/426 [=====] - 3s 7ms/step - loss: 1.0182 -
val_loss: 0.9411
Epoch 28/50

426/426 [=====] - 3s 7ms/step - loss: 1.0182 -
val_loss: 0.9409
Epoch 29/50
426/426 [=====] - 3s 6ms/step - loss: 1.0183 -
val_loss: 0.9409
Epoch 30/50
426/426 [=====] - 3s 6ms/step - loss: 1.0183 -
val_loss: 0.9407
Epoch 31/50
426/426 [=====] - 3s 6ms/step - loss: 1.0181 -
val_loss: 0.9417
Epoch 32/50
426/426 [=====] - 4s 9ms/step - loss: 1.0182 -
val_loss: 0.9407
Epoch 33/50
426/426 [=====] - 3s 6ms/step - loss: 1.0182 -
val_loss: 0.9408
Epoch 34/50
426/426 [=====] - 3s 6ms/step - loss: 1.0180 -
val_loss: 0.9410
Epoch 35/50
426/426 [=====] - 3s 6ms/step - loss: 1.0181 -
val_loss: 0.9407
Epoch 36/50
426/426 [=====] - 3s 8ms/step - loss: 1.0180 -
val_loss: 0.9408
Epoch 37/50
426/426 [=====] - 3s 8ms/step - loss: 1.0180 -
val_loss: 0.9407
Epoch 38/50
426/426 [=====] - 3s 6ms/step - loss: 1.0179 -
val_loss: 0.9413
Epoch 39/50
426/426 [=====] - 2s 6ms/step - loss: 1.0177 -
val_loss: 0.9415
Epoch 40/50
426/426 [=====] - 2s 6ms/step - loss: 1.0178 -
val_loss: 0.9407
Epoch 41/50
426/426 [=====] - 3s 8ms/step - loss: 1.0177 -
val_loss: 0.9406
Epoch 42/50
426/426 [=====] - 3s 6ms/step - loss: 1.0177 -
val_loss: 0.9411
Epoch 43/50
426/426 [=====] - 2s 6ms/step - loss: 1.0174 -
val_loss: 0.9411
Epoch 44/50

```

426/426 [=====] - 3s 6ms/step - loss: 1.0175 -
val_loss: 0.9411
Epoch 45/50
426/426 [=====] - 3s 7ms/step - loss: 1.0175 -
val_loss: 0.9408
Epoch 46/50
426/426 [=====] - 4s 9ms/step - loss: 1.0173 -
val_loss: 0.9416
Epoch 47/50
426/426 [=====] - 3s 6ms/step - loss: 1.0172 -
val_loss: 0.9410
Epoch 48/50
426/426 [=====] - 3s 6ms/step - loss: 1.0170 -
val_loss: 0.9420
Epoch 49/50
426/426 [=====] - 2s 6ms/step - loss: 1.0167 -
val_loss: 0.9410
Epoch 50/50
426/426 [=====] - 4s 9ms/step - loss: 1.0169 -
val_loss: 0.9411

```

[173]: <keras.callbacks.History at 0x7bd561e87b80>

[174]: `model.summary()`

Model: "sequential_2"

Layer (type)	Output Shape	Param #
lstm_10 (LSTM)	(None, 4)	112
dense_1 (Dense)	(None, 1)	5

```

=====
Total params: 117
Trainable params: 117
Non-trainable params: 0
=====

```

[175]: `#invert pridiction`
`trainprediction=model.predict(xtrain)`
`testprediction=model.predict(xtest)`

```

213/213 [=====] - 1s 2ms/step
71/71 [=====] - 0s 4ms/step

```

[156]: `from sklearn.metrics import mean_squared_error`

```
[176]: print(mean_squared_error(trainprediction,ytrain))
```

```
1.0163514446549322
```

```
[177]: ytrain
```


```
[177]: array([-0.28310702, -0.29424361,  0.24985899, ..., -0.27297519,  
          -0.01093009, -0.52144498])
```

```
[178]: trainprediction
```

```
[178]: array([[ 0.04358709],  
          [-0.05365378],  
          [-0.02182905],  
          ...,  
          [ 0.01246042],  
          [-0.00134046],  
          [-0.00943948]], dtype=float32)
```

```
[180]: plt.plot(trainprediction)
```

```
[180]: [<matplotlib.lines.Line2D at 0x7bd56a2ab4c0>]
```



[]: